

nasa-ptml

July 3, 2022

1 Introduction

Authors:

- Baptiste Bourdet
- Philippe Bernet
- Marius Dubosc
- Hugo Levy

The dataset comes from Kaggle: <https://www.kaggle.com/datasets/sameepvani/nasa-nearest-earth-objects>.

In this report we will try to analyze this data and compute models of both supervised and unsupervised learning to respond to a problem that was highlighted in many movies of science fiction: are any objects currently in orbit a danger to either satellites or earth.

1.1 Loading the data

First we need to load the data and analyse the different component it is made out of.

| | id | name | est_diameter_min | est_diameter_max | \ |
|-------|----------|---------------------|------------------|------------------|---|
| 0 | 2162635 | 162635 (2000 SS164) | 1.198271 | 2.679415 | |
| 1 | 2277475 | 277475 (2005 WK4) | 0.265800 | 0.594347 | |
| 2 | 2512244 | 512244 (2015 YE18) | 0.722030 | 1.614507 | |
| 3 | 3596030 | (2012 BV13) | 0.096506 | 0.215794 | |
| 4 | 3667127 | (2014 GE35) | 0.255009 | 0.570217 | |
| ... | ... | ... | ... | ... | |
| 90831 | 3763337 | (2016 VX1) | 0.026580 | 0.059435 | |
| 90832 | 3837603 | (2019 AD3) | 0.016771 | 0.037501 | |
| 90833 | 54017201 | (2020 JP3) | 0.031956 | 0.071456 | |
| 90834 | 54115824 | (2021 CN5) | 0.007321 | 0.016370 | |
| 90835 | 54205447 | (2021 TW7) | 0.039862 | 0.089133 | |

| | relative_velocity | miss_distance | orbiting_body | sentry_object | \ |
|---|-------------------|---------------|---------------|---------------|---|
| 0 | 13569.249224 | 5.483974e+07 | Earth | False | |
| 1 | 73588.726663 | 6.143813e+07 | Earth | False | |
| 2 | 114258.692129 | 4.979872e+07 | Earth | False | |
| 3 | 24764.303138 | 2.543497e+07 | Earth | False | |

| | | | | |
|-------|--------------|--------------|-------|-------|
| 4 | 42737.733765 | 4.627557e+07 | Earth | False |
| ... | ... | ... | ... | ... |
| 90831 | 52078.886692 | 1.230039e+07 | Earth | False |
| 90832 | 46114.605073 | 5.432121e+07 | Earth | False |
| 90833 | 7566.807732 | 2.840077e+07 | Earth | False |
| 90834 | 69199.154484 | 6.869206e+07 | Earth | False |
| 90835 | 27024.455553 | 5.977213e+07 | Earth | False |

| | absolute_magnitude | hazardous |
|-------|--------------------|-----------|
| 0 | 16.73 | False |
| 1 | 20.00 | True |
| 2 | 17.83 | False |
| 3 | 22.20 | False |
| 4 | 20.09 | True |
| ... | ... | ... |
| 90831 | 25.00 | False |
| 90832 | 26.00 | False |
| 90833 | 24.60 | False |
| 90834 | 27.80 | False |
| 90835 | 24.12 | False |

[90836 rows x 10 columns]

The data is composed of the following columns:

- id: index number
- name: name of the object
- est_dimater_min: smallest size of the object in km
- est_dimater_max: biggest size of the object in km
- relative_velocity: velocity relative to Earth in km/h
- orbiting_body: the body the object is orbiting (Earth, Sun, the Moon ...)
- sentry_object: whether or not the object is tracked by the sentry system of the nasa
- absolute_magnitude: visibility index, the smaller it is, the brighther the object it, the magnitude of the sun is -27 for example
- **hazardous**: whether or not the object is considerer a potential threat by the nasa, it is this column we will want to monitor

2 Analysis

2.1 Basic statistics

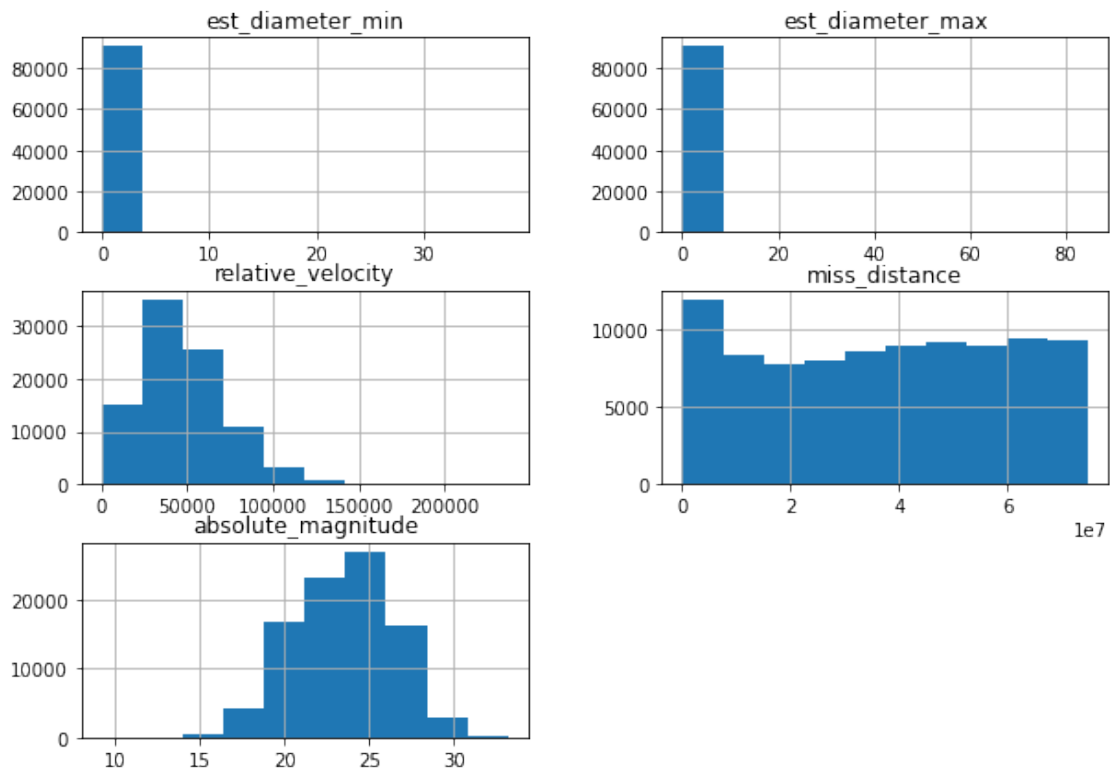
First we will look at statistics on the dataset.

Length of the dataset is 90836

Summary of all numerical values :

| | est_diameter_min | est_diameter_max | relative_velocity | miss_distance \ |
|-------|------------------|------------------|-------------------|-----------------|
| count | 90836.000000 | 90836.000000 | 90836.000000 | 9.083600e+04 |
| mean | 0.127432 | 0.284947 | 48066.918918 | 3.706655e+07 |
| std | 0.298511 | 0.667491 | 25293.296961 | 2.235204e+07 |
| min | 0.000609 | 0.001362 | 203.346433 | 6.745533e+03 |
| 25% | 0.019256 | 0.043057 | 28619.020645 | 1.721082e+07 |
| 50% | 0.048368 | 0.108153 | 44190.117890 | 3.784658e+07 |
| 75% | 0.143402 | 0.320656 | 62923.604633 | 5.654900e+07 |
| max | 37.892650 | 84.730541 | 236990.128088 | 7.479865e+07 |

| | absolute_magnitude |
|-------|--------------------|
| count | 90836.000000 |
| mean | 23.527103 |
| std | 2.894086 |
| min | 9.230000 |
| 25% | 21.340000 |
| 50% | 23.700000 |
| 75% | 25.700000 |
| max | 33.200000 |



<class 'pandas.core.frame.DataFrame'>
RangeIndex: 90836 entries, 0 to 90835

Data columns (total 10 columns):

| # | Column | Non-Null Count | Dtype |
|---|--------------------|----------------|----------|
| 0 | id | 90836 non-null | object |
| 1 | name | 90836 non-null | object |
| 2 | est_diameter_min | 90836 non-null | float64 |
| 3 | est_diameter_max | 90836 non-null | float64 |
| 4 | relative_velocity | 90836 non-null | float64 |
| 5 | miss_distance | 90836 non-null | float64 |
| 6 | orbiting_body | 90836 non-null | category |
| 7 | sentry_object | 90836 non-null | bool |
| 8 | absolute_magnitude | 90836 non-null | float64 |
| 9 | hazardous | 90836 non-null | bool |

dtypes: bool(2), category(1), float64(5), object(2)
memory usage: 5.1+ MB


Visualization of categorical values :

| | orbiting_body | sentry_object | hazardous |
|--------|---------------|---------------|-----------|
| count | 90836 | 90836 | 90836 |
| unique | 1 | 1 | 2 |
| top | Earth | False | False |
| freq | 90836 | 90836 | 81996 |

As we can see we have only one **orbiting body** and only one value for the **sentry**. Those two columns can therefore be removed from the dataset safely. The final column, the one we want to predict with the models, seems to have 10% of dangerous objects, the outliers we will try to identify.


2.2 Correlation

The correlation matrix of our data is the following:



We can see in the matrix that the magnitude is negatively correlated with most of the other variables. This implies that objects with a high magnitude, meaning not very visible objects, are usually smaller and slower. They also more importantly do not constitute a threat seeing how the magnitude is negatively correlated with the hazardous.


When computing the correlation matrix only for the hazardous objects, we can see an even greater correlation between the magnitude and the size.



2.3 3D representation


We can also visualize multiple columns at once to see patterns.

Hazardous and Non Hazardous objects in Earth's vicinity



As can be seen on this first representation, the velocity seems to be slightly higher for objects considered a threat compared to the rest. The magnitude also seems to be capped at 20 for the hazardous objects, a high magnitude implying a very dim object in the darkness of space.

Hazardous and Non Hazardous objects in Earth's vicinity



On this second graph we can further see that the velocity seems to be an important factor but on the other end the miss distance is not very representative. Indeed the miss distance is not that important considering objects could be shifted out of their orbit very easily by the cosmic billiard played in the solar system by gravity. An object that missed earth by a lot could still be a threat.

3 Supervised Learning

First, we clean the dataset to get quantitative data to determine whether or not the object is considered a potential threat by NASA, it is this column we will want to monitor. We remove constant values (`sentry_object=False` and `orbiting_body=Earth`)

```
/opt/conda/lib/python3.9/site-packages/pandas/core/indexes/base.py:6982:
FutureWarning: In a future version, the Index constructor will not infer numeric
dtypes when passed object-dtype sequences (matching Series behavior)
    return Index(sequences[0], name=names)
```


| | est_diameter_min | est_diameter_max | relative_velocity \ |
|----------|------------------|------------------|---------------------|
| id | | | |
| 2162635 | 1.198271 | 2.679415 | 13569.249224 |
| 2277475 | 0.265800 | 0.594347 | 73588.726663 |
| 2512244 | 0.722030 | 1.614507 | 114258.692129 |
| 3596030 | 0.096506 | 0.215794 | 24764.303138 |
| 3667127 | 0.255009 | 0.570217 | 42737.733765 |
| ... | ... | ... | ... |
| 3763337 | 0.026580 | 0.059435 | 52078.886692 |
| 3837603 | 0.016771 | 0.037501 | 46114.605073 |
| 54017201 | 0.031956 | 0.071456 | 7566.807732 |
| 54115824 | 0.007321 | 0.016370 | 69199.154484 |
| 54205447 | 0.039862 | 0.089133 | 27024.455553 |

| | miss_distance | absolute_magnitude | hazardous |
|----------|---------------|--------------------|-----------|
| id | | | |
| 2162635 | 5.483974e+07 | 16.73 | False |
| 2277475 | 6.143813e+07 | 20.00 | True |
| 2512244 | 4.979872e+07 | 17.83 | False |
| 3596030 | 2.543497e+07 | 22.20 | False |
| 3667127 | 4.627557e+07 | 20.09 | True |
| ... | ... | ... | ... |
| 3763337 | 1.230039e+07 | 25.00 | False |
| 3837603 | 5.432121e+07 | 26.00 | False |
| 54017201 | 2.840077e+07 | 24.60 | False |
| 54115824 | 6.869206e+07 | 27.80 | False |
| 54205447 | 5.977213e+07 | 24.12 | False |

[90836 rows x 6 columns]

3.1 Initialize train and test sets

3.2 Let's try multiple models

```
===== LogisticRegressionCV(class_weight='balanced') =====
```

```
----- Classification Report -----
```

train :

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.94 | 0.38 | 0.54 | 54969 |
| True | 0.12 | 0.76 | 0.20 | 5891 |
| accuracy | | | 0.41 | 60860 |
| macro avg | 0.53 | 0.57 | 0.37 | 60860 |
| weighted avg | 0.86 | 0.41 | 0.50 | 60860 |

test :

| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|------|------|------|-------|
| False | 0.93 | 0.38 | 0.54 | 27027 |
| True | 0.12 | 0.76 | 0.20 | 2949 |
| accuracy | | | 0.42 | 29976 |
| macro avg | 0.53 | 0.57 | 0.37 | 29976 |
| weighted avg | 0.85 | 0.42 | 0.51 | 29976 |

----- Accuracy Score -----

train :
0.41390075583305946
test :
0.4150987456631972

----- Confusion matrix on test -----

```
[[0.37758538 0.62241462]
 [0.24109868 0.75890132]]
```

===== KNeighborsClassifier() =====

----- Classification Report -----

train :

| | | | | |
|--------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| False | 0.91 | 0.99 | 0.95 | 54969 |
| True | 0.66 | 0.11 | 0.18 | 5891 |
| accuracy | | | 0.91 | 60860 |
| macro avg | 0.79 | 0.55 | 0.57 | 60860 |
| weighted avg | 0.89 | 0.91 | 0.88 | 60860 |

test :

| | | | | |
|--------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| False | 0.90 | 0.99 | 0.94 | 27027 |
| True | 0.21 | 0.03 | 0.06 | 2949 |
| accuracy | | | 0.89 | 29976 |
| macro avg | 0.56 | 0.51 | 0.50 | 29976 |
| weighted avg | 0.83 | 0.89 | 0.86 | 29976 |

----- Accuracy Score -----

train :
0.9083470259612225
test :
0.8922471310381639

----- Confusion matrix on test -----

```
[[0.98590299 0.01409701]
```

```
[0.9660902  0.0339098  ]
```

```
===== DecisionTreeClassifier(random_state=0) =====
```

```
----- Classification Report -----
```

```
train :
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 1.00 | 1.00 | 1.00 | 54969 |
| True | 1.00 | 1.00 | 1.00 | 5891 |
| accuracy | | | 1.00 | 60860 |
| macro avg | 1.00 | 1.00 | 1.00 | 60860 |
| weighted avg | 1.00 | 1.00 | 1.00 | 60860 |

```
test :
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.94 | 0.94 | 0.94 | 27027 |
| True | 0.44 | 0.45 | 0.44 | 2949 |
| accuracy | | | 0.89 | 29976 |
| macro avg | 0.69 | 0.69 | 0.69 | 29976 |
| weighted avg | 0.89 | 0.89 | 0.89 | 29976 |

```
----- Accuracy Score -----
```

```
train :
```

```
1.0
```

```
test :
```

```
0.8899786495863358
```

```
----- Confusion matrix on test -----
```

```
[[0.93817294 0.06182706]  
 [0.55171244 0.44828756]]
```

```
===== RandomForestClassifier(max_depth=3, n_estimators=50,
```

```
random_state=0) =====
```

```
----- Classification Report -----
```

```
train :
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.91 | 1.00 | 0.95 | 54969 |
| True | 0.87 | 0.13 | 0.22 | 5891 |
| accuracy | | | 0.91 | 60860 |
| macro avg | 0.89 | 0.56 | 0.59 | 60860 |
| weighted avg | 0.91 | 0.91 | 0.88 | 60860 |

```
test :
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.91 | 1.00 | 0.95 | 27027 |
| True | 0.85 | 0.11 | 0.20 | 2949 |
| accuracy | | | 0.91 | 29976 |
| macro avg | 0.88 | 0.56 | 0.58 | 29976 |
| weighted avg | 0.91 | 0.91 | 0.88 | 29976 |

----- Accuracy Score -----

train :

0.9136214262241209

test :


0.9107619428876434


----- Confusion matrix on test -----

[[0.997817 0.002183]

[0.88708037 0.11291963]]

3.2.1 Dimension Reduction to facilitate classification





```
[5.31251159e-01 2.35392401e-01 1.33986946e-01 9.93694937e-02
 1.13445612e-21]
```

Score on training set : 1.0

Score on test set : 0.8840405657859621


Score on training set : 0.9060138021689123

Score on test set : 0.9042233787029623

```
[[25247 1780]
 [ 1696 1253]]
```

False negative : 1696

False positive : 1780



```
[[27018    9]
 [ 2862   87]]
False negative : 2862
False positive : 9
```

